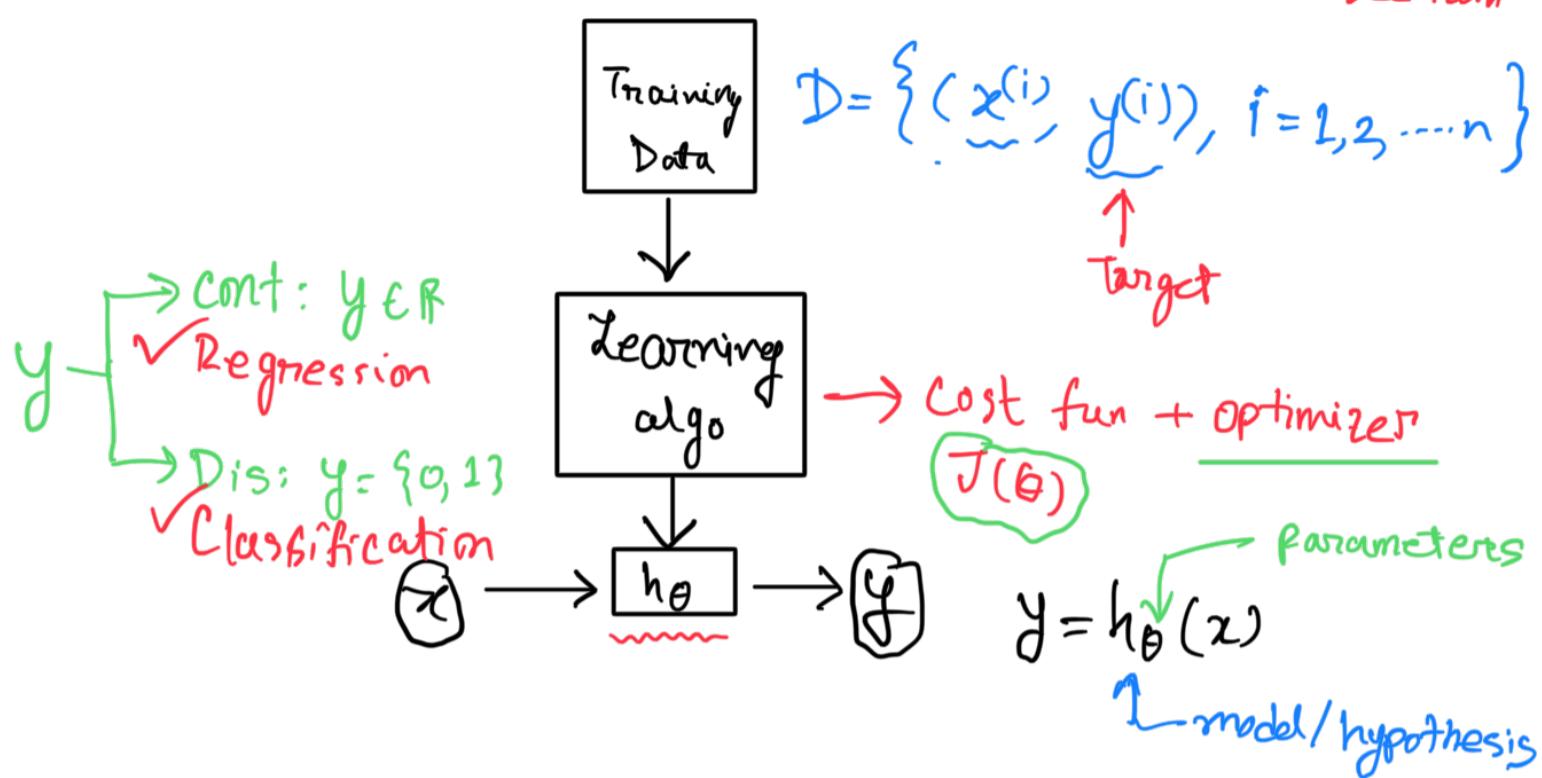


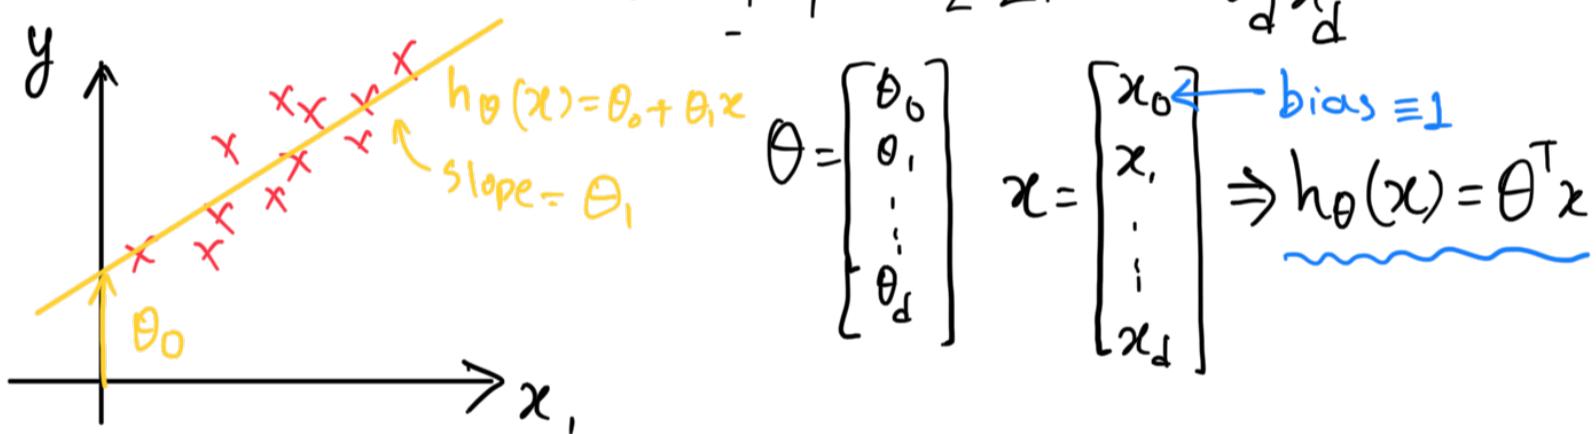
Supervised Learning

$$\mathbf{x}^{(i)} = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \end{bmatrix} \begin{array}{l} \xleftarrow{1} \\ \xleftarrow{\text{sq. ft}} \\ \xleftarrow{\# \text{bedrooms}} \end{array}$$



Linear Regression \rightarrow Solving regression problem

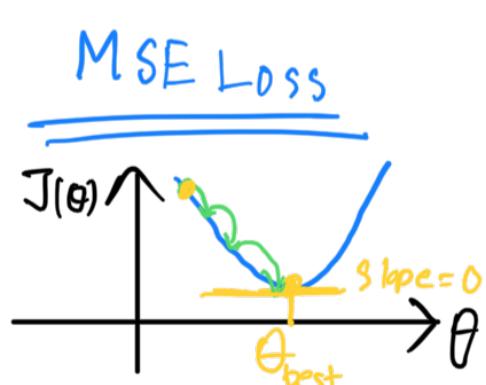
$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_d x_d$$



$$\text{Cost function: } J_{\theta}(x) = \frac{1}{n} \sum_{i=1}^n J_i(\theta)$$

Total sample cost

$$\begin{aligned} J_i(\theta) &= [h_{\theta}(x^i) - y^i]^2 \\ &= [\theta^T x^i - y^i]^2 \end{aligned}$$



Optimizer; Gradient Descent

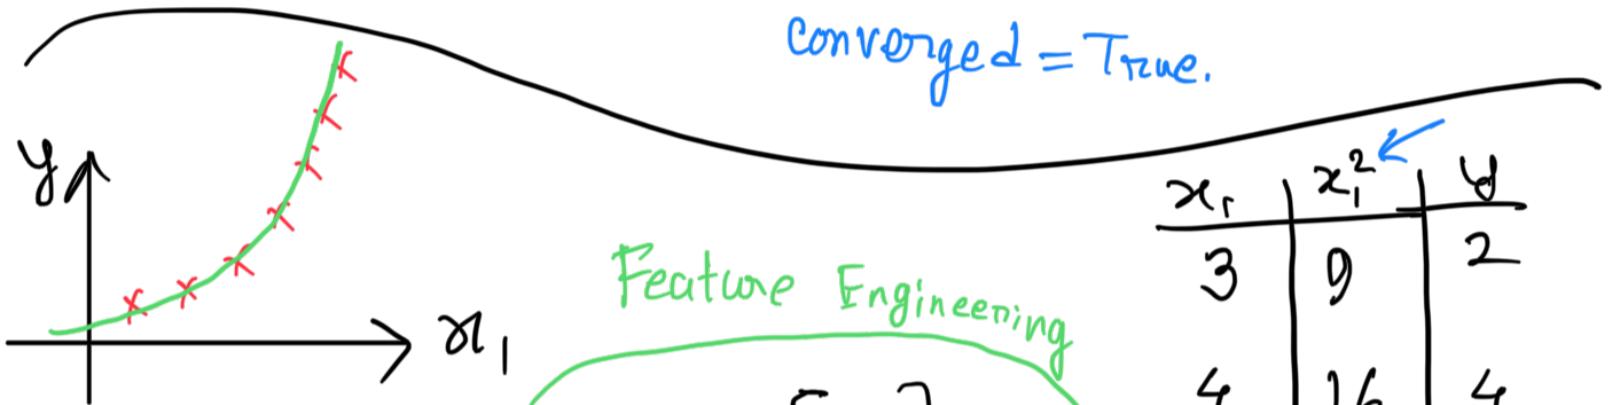
$\theta^0 \leftarrow$ iteration no
 $\theta = \text{random}$

while not converge?:

$$\theta^{k+1} = \theta^k - \alpha \nabla_{\theta} J \rightarrow \text{gradient of } J \text{ w.r.t. } \theta$$

if $\nabla_{\theta} J \approx 0$: step size

Converged = True.



$$x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix}$$

Feature Engineering

$$\phi(x) = \begin{bmatrix} x_0 \\ x_1 \\ x_1^2 \end{bmatrix} = x'$$

x_0	x_1^2	y
3	9	2
4	16	4
5	25	16

$$h_{\theta}(x) = \theta^T x' = \theta_0 + \theta_1 x_1 + \theta_2 x_1^2$$

$$J(\theta) = \frac{1}{n} \sum_{i=0}^n \left[\theta^T x^{(i)} - y^{(i)} \right]^2$$

$$= \frac{1}{n} \sum_{i=0}^n \left[\theta_0 + \theta_1 x_1 + \theta_2 x_1^2 - y^{(i)} \right]^2$$

Classification Problem

y is discrete.

e.g. Binary $\rightarrow y = \{0, 1\}$

x_1	x_2	x_3	y
3	0.2	24	0
6	-0.5	50	1

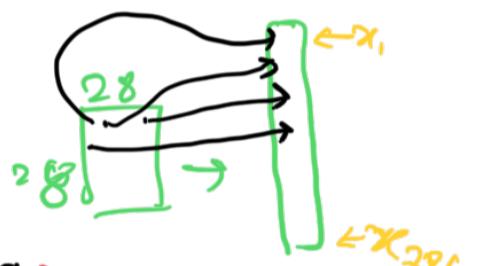
← size density Age ← benign / malignant

Example

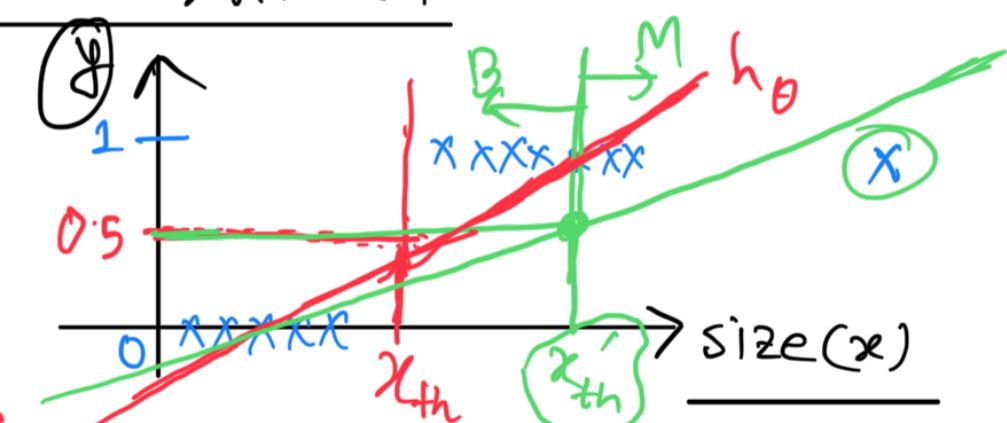
image pixels \rightarrow cat / dog
 $x = [x_0, x_1, \dots, x_d]$ $y = \{0, 1\}$

tumor size \rightarrow benign / malignant
 $x = [x_0, x_1]$ $y = \{0, 1\}$

Email text \rightarrow spam / not spam



Tumor classification



$$h_\theta(x) = \theta^T x = \theta_0 + \theta_1 x$$

$$\begin{aligned} x > x_{th} &\equiv h_\theta(x) > 0.5 \Rightarrow \text{malignant} \\ x \leq x_{th} &\equiv h_\theta(x) < 0.5 \Rightarrow \text{benign} \end{aligned}$$

Classification problem \Rightarrow linear regression

① Sensitive to outliers

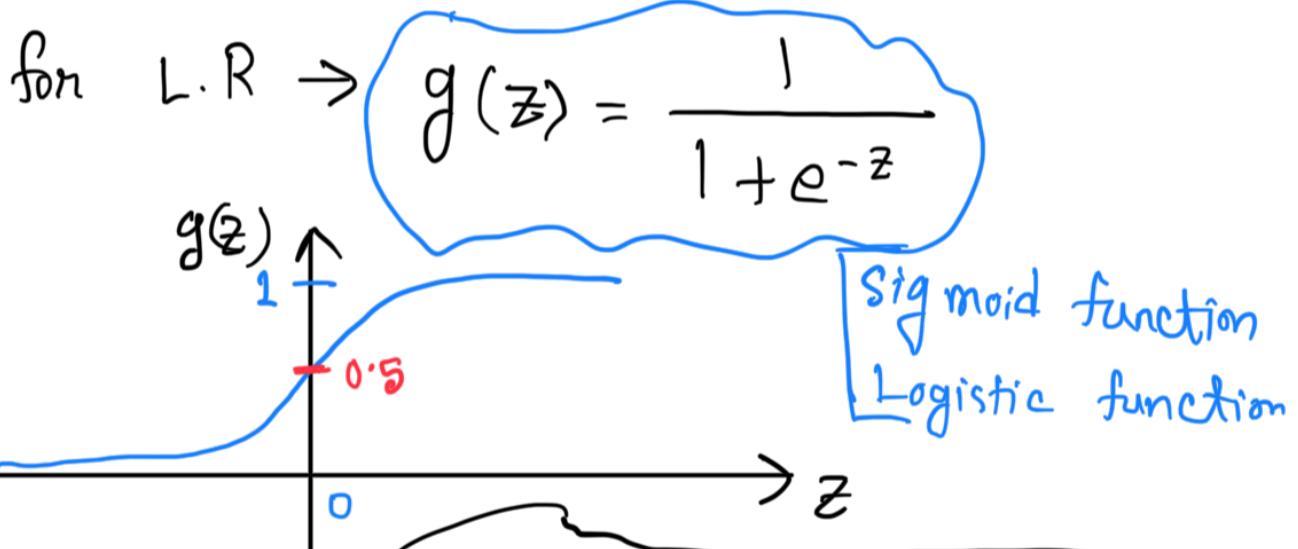
② $y = 0 \text{ or } 1$

$$h_{\theta}(x) = \theta^T x < 0 \text{ and } > 1$$

Logistic Regression

Idea: Transform $\theta^T x$ so that $0 \leq h_{\theta}(x) \leq 1$

$$h_{\theta}(x) = g(\underbrace{\theta^T x}_z) \text{ model}$$



$$\hat{y} = \begin{cases} 1 & \text{if } \frac{1}{1 + e^{-\theta^T x}} \geq 0.5 \\ 0 & \text{if } \frac{1}{1 + e^{-\theta^T x}} < 0.5 \end{cases}$$

$$y = \begin{cases} 1 & \text{if } \frac{1}{1 + e^{-\theta^T x}} \geq 0.5 \\ 0 & \text{if } \frac{1}{1 + e^{-\theta^T x}} < 0.5 \end{cases}$$

interpret: $h_{\theta}(x) = P(y=1 | x, \theta)$

probability that $y=1$

eg $x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ 3 \end{bmatrix}$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix} = \begin{bmatrix} 0.54 \\ 0.1 \end{bmatrix}$$

$$\theta^T x = \theta_0 x_0 + \theta_1 x_1 = 0.54 \times 1 + 0.1 \times 3 \\ = 0.84$$

$$g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}} = \frac{1}{1 + e^{-0.84}} = 0.2$$

$h_{\theta}(x) = g(\theta^T x) = 0.2$

$P(y=1 | x, \theta)$

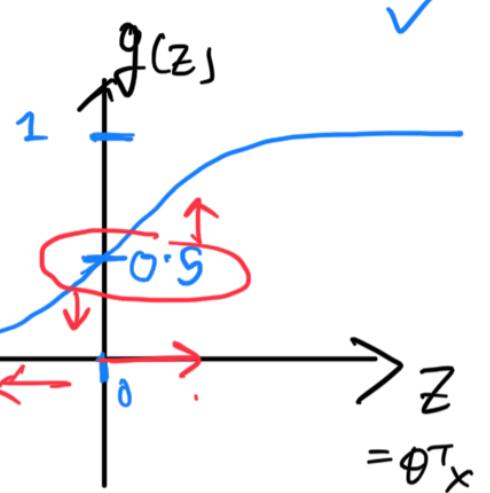
\hookrightarrow 20% probability that
the $y=1$ (tumor is
malignant)

Decision Boundary

$$h_{\theta}(x) = g(\theta^T x)$$

sigmoid function

$$= \frac{1}{1 + e^{-\theta^T x}}$$

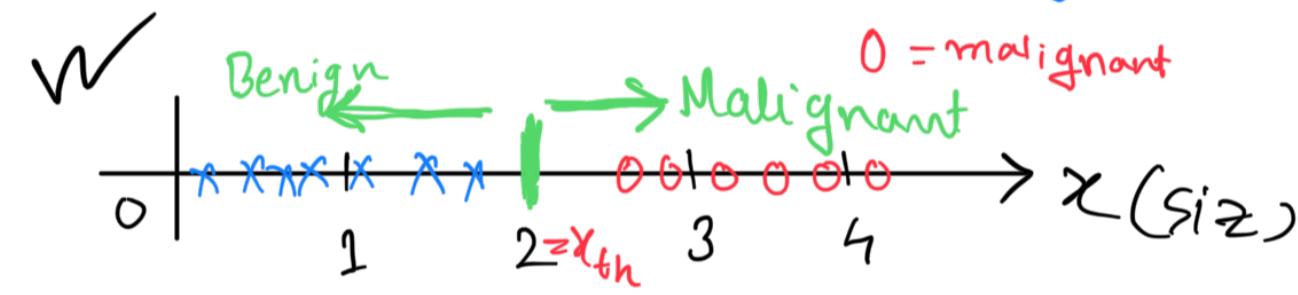


$$y = \begin{cases} 1 & \text{if } h_{\theta}(\theta^T x) > 0.5 \\ 0 & \text{else} \end{cases}$$

↳ when? $\Rightarrow z > 0$

if $z < 0 \Rightarrow \theta^T x > 0$

Eq. 1



$$h_{\theta}(x) = g(\theta^T x) = g(\theta_0 + \theta_1 x)$$

$$\rightarrow \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix} = \begin{bmatrix} -2 \\ 1 \end{bmatrix}$$

∴ $\theta^T x = \theta_0 + \theta_1 x = -2 + x$

∴ Decision boundary:

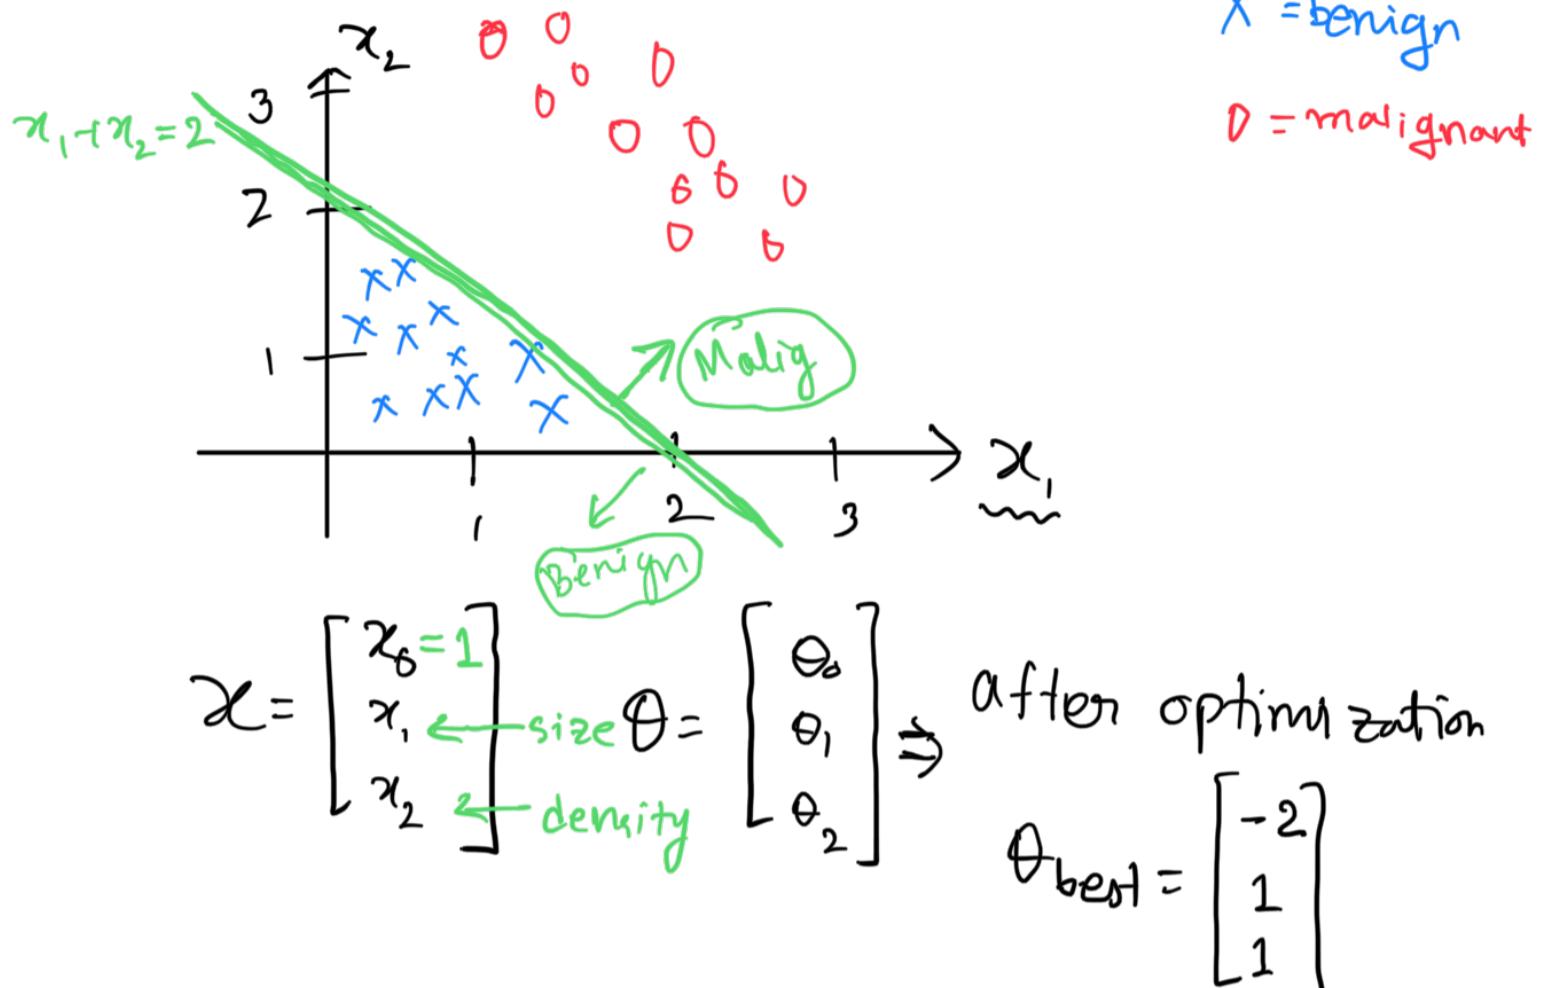
$$\rightarrow x = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ x \end{bmatrix}$$

$\theta^T x > 0$

$\Rightarrow -2 + x > 0$

$\Rightarrow \boxed{x > 2}$

eq-2



\therefore Decision boundary $\hat{\theta}^T x \geq 0$

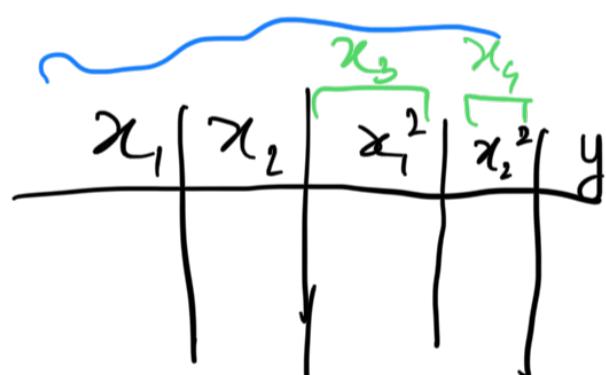
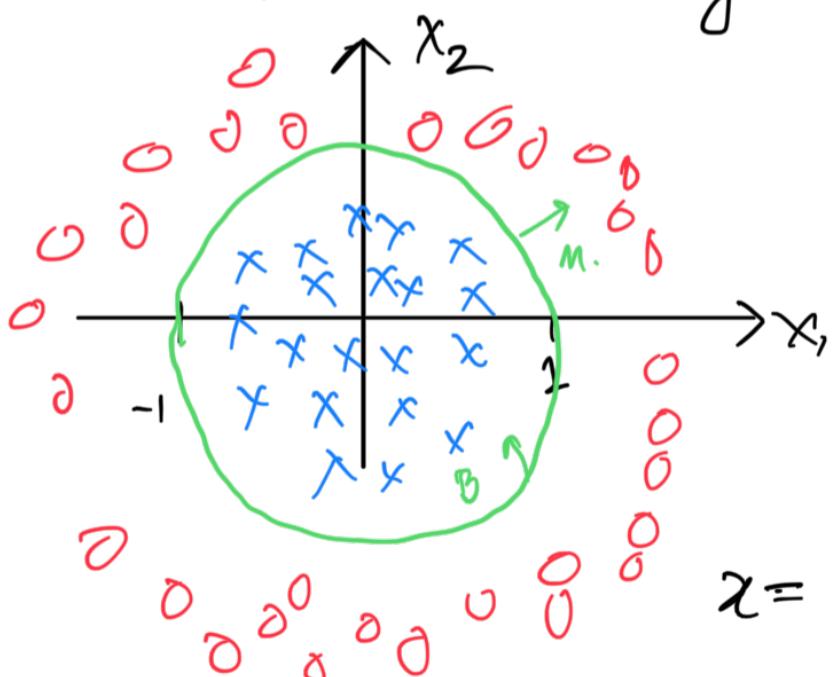
$$\Rightarrow \theta_0 + \theta_1 x_1 + \theta_2 x_2 \geq 0$$

$$x_1 + x_2 = 2$$

$$\Rightarrow -2 + x_1 + x_2 \geq 0$$

$$\Rightarrow x_1 + x_2 \geq 2$$

Eq-3 Complex boundary



$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_1^2 \\ x_2^2 \end{bmatrix}$$

$$\theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \theta_3 \\ \theta_4 \end{bmatrix}$$

$$\hat{\theta}^T x = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2 + \theta_4 x_2^2$$

$$\theta_{\text{best}} = \begin{bmatrix} -1 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

- Decision boundary

$$\Rightarrow -1 + 0 + 0 + x_1^2 + x_2^2 \geq 0$$

$$x_1^2 + x_2^2 = 1$$

$$\Rightarrow x_1^2 + x_2^2 \geq 1$$

$$h_{\theta}(x) = g(\theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_1^2$$

$$+ \theta_4 x_2^2 + \theta_5 x_1^3 x_2 + \theta_6 x_1^2 x_2^2)$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \\ x_4 \\ x_5 \\ x_6 \\ x_7 \end{bmatrix} \Rightarrow \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \\ \theta_3 \\ \theta_4 \\ \theta_5 \\ \theta_6 \\ \theta_7 \end{bmatrix} + \theta_8 \sin(x_1))$$

$$h_{\theta} = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

Fitting the Model \equiv Finding Best θ [Learning Algo]

$$D = \{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(n)}, y^{(n)})\}$$

$$\rightarrow x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_d \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_d \end{bmatrix}$$

$y = \{0, 1\}$
 n -sample
 $g(z) = \text{sigmoid}$

$$\rightarrow h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

Task: Find the "best" θ → θ that performs best of D → minimizes some cost function $J(\theta)$

Cost function

$$J(\theta) = \frac{1}{n} \sum_{i=1}^n J_i(\theta)$$

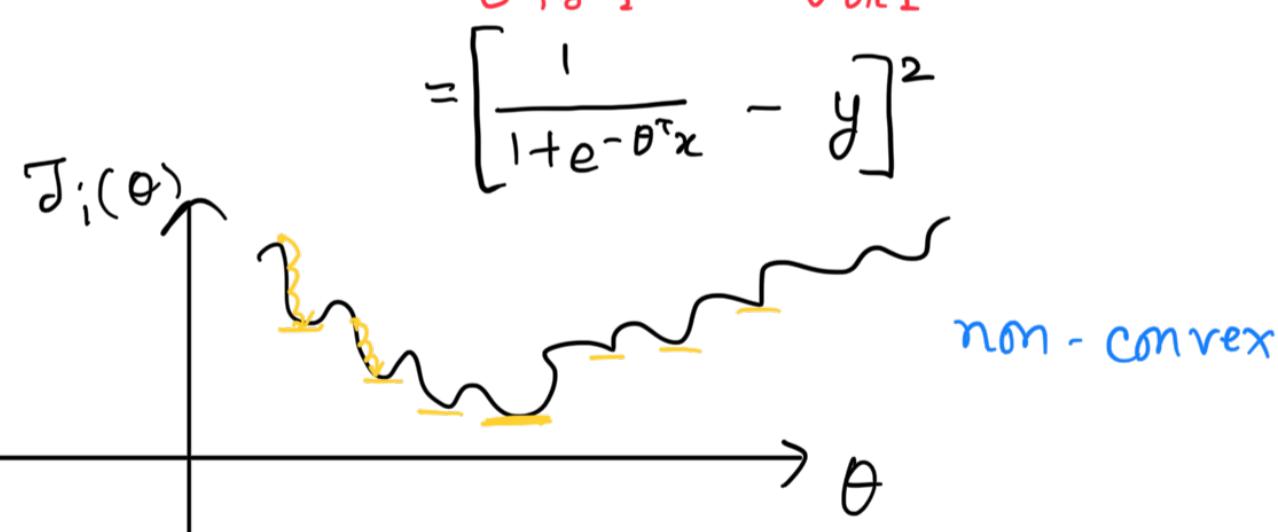
Cost for the i th training sample

Opt 1: MSE

$$J_i(\theta) = [h_{\theta}(x) - y]^2$$

$0 \text{ or } 1$

$x = x^{(i)}$
 $y = y^{(i)}$



Opt -2 : Binary Cross-Entropy Loss

log base=e

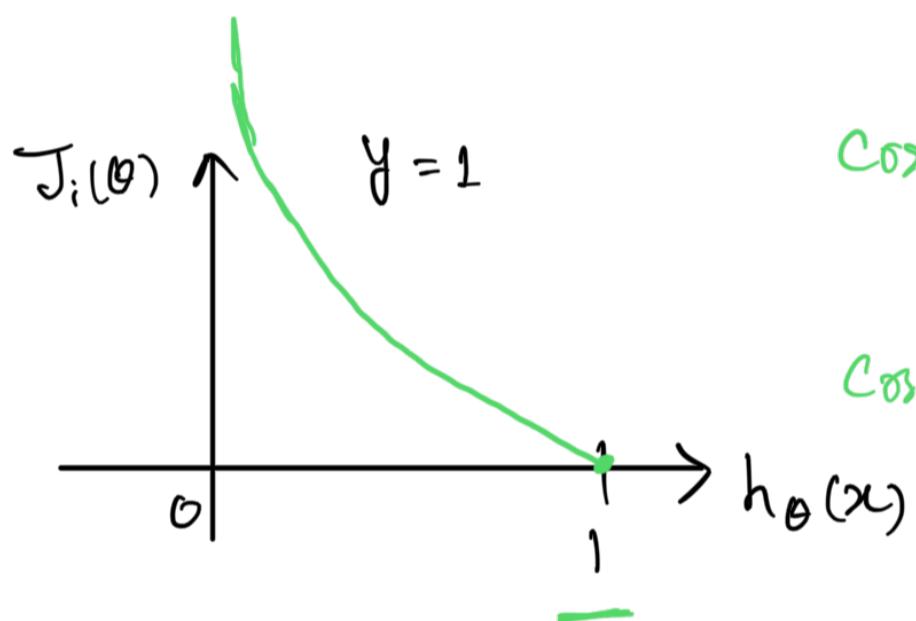
cost function
for ith sample

output of model

$$J_i(\theta) = \begin{cases} -\log(h_\theta(x)) & \text{if } y=1 \\ -\log(1-h_\theta(x)) & \text{if } y=0 \end{cases}$$

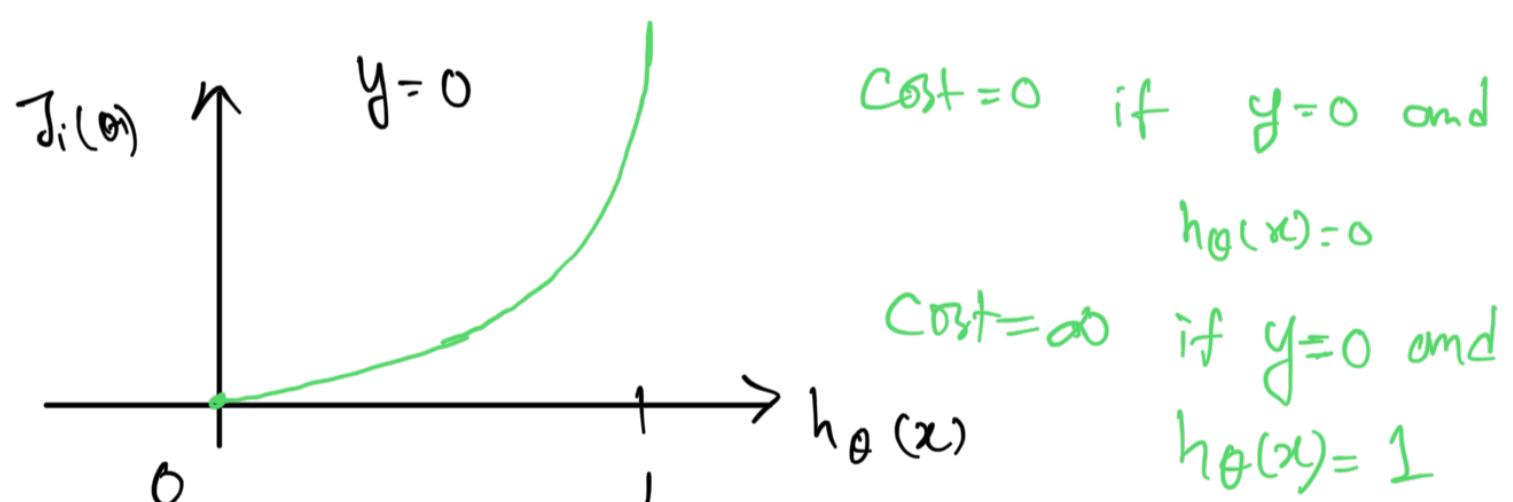
$$\Rightarrow J_i(\theta) = -[y \times \log(h_\theta(x)) + (1-y) \log(1-h_\theta(x))]$$

zero if $y=0$ zero if $y=1$



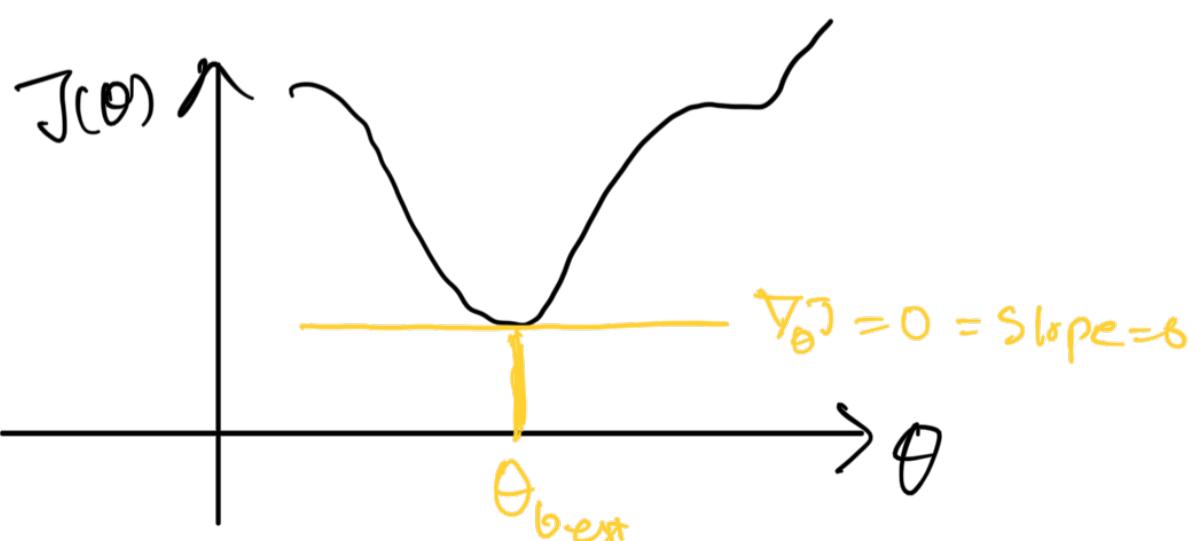
Cost = 0 if $y=1$ and $h_\theta(x)=1$

Cost = ∞ if $y=1$ and $h_\theta(x)=0$



Cost = 0 if $y=0$ and $h_\theta(x)=0$

Cost = ∞ if $y=0$ and $h_\theta(x)=1$



Minimizing the Cost

Optimizer

① Gradient descent

$\theta^0 = \text{random}$ ↪

while not converge? ↪

$$\theta^{k+1} = \theta^k - \alpha \nabla_{\theta} J$$

if $\nabla_{\theta} J \approx 0$:

converged = True.

② Adam

③ Ada boost

④ LIBSFS

Complex But

faster convergence

Logistic Regression Summary

$$D = \{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(n)}, y^{(n)})\}$$

$$x = \begin{bmatrix} x_0 \\ x_1 \\ \vdots \\ x_d \end{bmatrix} \quad \theta = \begin{bmatrix} \theta_0 \\ \theta_1 \\ \vdots \\ \theta_d \end{bmatrix}$$

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

$$J(\theta) = -\frac{1}{n} \sum_{i=1}^n [y^{(i)} \log(h_{\theta}(x^{(i)})) + (1-y^{(i)}) \log(1-h_{\theta}(x^{(i)}))]$$

↑ Binary cross entropy loss

$$\theta_{\text{best}} = \arg \min J(\theta)$$

Optimizer → G.D.
→ S.G.D
→ Adam

