

Constraint Satisfaction

Russell & Norvig Ch. 6.1-6.4

Informal Definition of CSP


- CSP = Constraint Satisfaction Problem
- Given
 - (1) a **finite set of variables**
 - (2) each with a domain of possible values (often finite)
 - (3) **a set of constraints that limit the values** the variables can take on
- A **solution** is an assignment of a value to each variable such that the **constraints are all satisfied**.
- Tasks might be to decide **if a solution exists**, to find a **solution**, to find all solutions, or to find the “**best solution**” according to some metric (objective function).

Today's Class

- Constraint Processing / Constraint Satisfaction Problem (CSP) paradigm
- Algorithms for CSPs
 - Backtracking (systematic search)
 - Constraint propagation (k-consistency)
 - Variable and value ordering heuristics
 - Intelligent backtracking

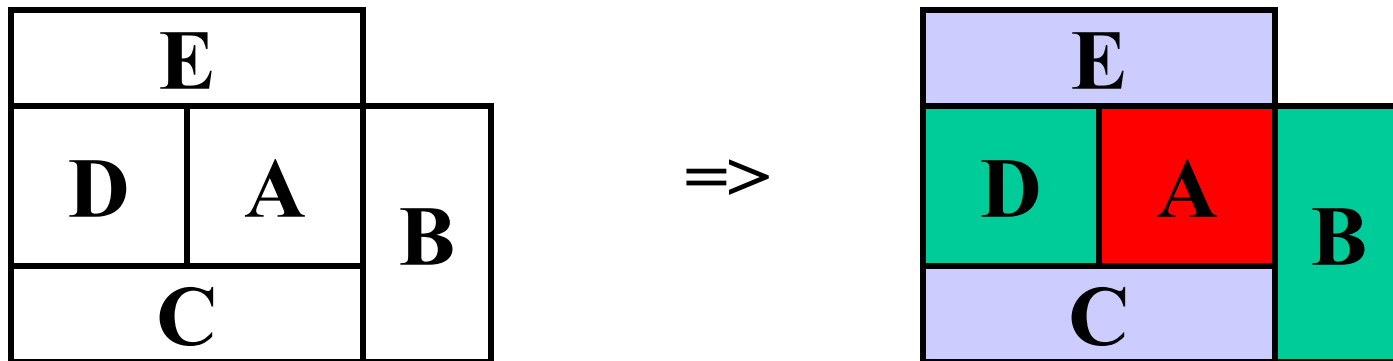
Informal Example: Map Coloring

- Color the following map using three colors (red, green, blue) such that no two adjacent regions have the same color.



Map Coloring II

- Variables: A, B, C, D, E all of domain RGB
- Domains: RGB = {red, green, blue}
- Constraints: $A \neq B$, $A \neq C$, $A \neq E$, $A \neq D$, $B \neq C$, $C \neq D$, $D \neq E$
- One solution: A=red, B=green, C=blue, D=green, E=blue



Formal Definition of a Constraint Network (CN)

A constraint network (CN) consists of

- a set of variables $X = \{x_1, x_2, \dots, x_n\}$
 - each with an associated domain of values $\{d_1, d_2, \dots, d_n\}$.
 - the domains are typically finite
- a set of constraints $\{c_1, c_2, \dots, c_m\}$ where
 - each constraint defines a predicate which is a relation over a particular subset of X .
 - e.g., C_i involves variables $\{X_{i1}, X_{i2}, \dots, X_{ik}\}$ and defines the relation $R_i \subseteq D_{i1} \times D_{i2} \times \dots \times D_{ik}$
- **Unary** constraint: only **involves one variable**
- **Binary** constraint: only **involves two variables**

Example (Class Scheduling)

- Given a list of courses to be taught, classrooms available, time slots, and professors who can teach certain courses, can classes be scheduled?
- Variables: Courses offered (C_1, \dots, C_i), classrooms (R_1, \dots, R_j), time (T_1, \dots, T_k).
- Domains:
 - $DC_i = \{\text{professors who can teach course } i\}$
 - $DR_j = \{\text{room numbers}\}$
 - $DT_k = \{\text{time slots}\}$
- Constraints:
 - Maximum 1 class per room in each time slot.
 - A professor cannot teach 2 classes in the same time slot.
 - A professor cannot teach more than 2 classes.

Typical Tasks for CSP

- Solutions:
 - Does a solution exist?
 - Find one solution
 - Find all solutions
 - Given a partial instantiation, do any of the above
- Transform the CN into an equivalent CN that is easier to solve.

Solving Constraint Problems

- Systematic search
 - Generate and test
 - Backtracking
- Constraint propagation (consistency)
- Variable ordering heuristics
- Value ordering heuristics
- Backjumping and dependency-directed backtracking

Systematic Search: Backtracking

(a.k.a. depth-first search!)

- Consider the variables in some order
- Pick an unassigned variable and give it a provisional value such that it is consistent with all of the constraints
- If no such assignment can be made, we've reached a dead end and need to backtrack to the previous variable
- Continue this process until a solution is found or we backtrack to the initial variable and have exhausted all possible values

Problems with Backtracking

- Thrashing: keep repeating the same failed variable assignments
 - Consistency checking can help
 - Intelligent backtracking schemes can also help
- Inefficiency: can explore areas of the search space that aren't likely to succeed
 - Variable ordering can help

Consistency


- Node consistency
 - A node X is **node-consistent** if every value in the domain of X is consistent with X 's unary constraints
 - A graph is node-consistent if all nodes are node-consistent
- Arc consistency
 - An arc (X, Y) is **arc-consistent** if, for every value x of X , there is a value y for Y that satisfies the constraint represented by the arc.
 - A graph is arc-consistent if all arcs are arc-consistent.
- To create arc consistency, we perform **constraint propagation**: that is, we repeatedly reduce the domain of each variable to be consistent with its arcs

Example: Map-Coloring



- **Variables** WA, NT, Q, NSW, V, SA, T
- **Domains** $D_i = \{\text{red, green, blue}\}$
- **Constraints**: adjacent regions must have different colors
e.g., $WA \neq NT$, or $(WA, NT) \in \{(\text{red, green}), (\text{red, blue}), (\text{green, red}), (\text{green, blue}), (\text{blue, red}), (\text{blue, green})\}$

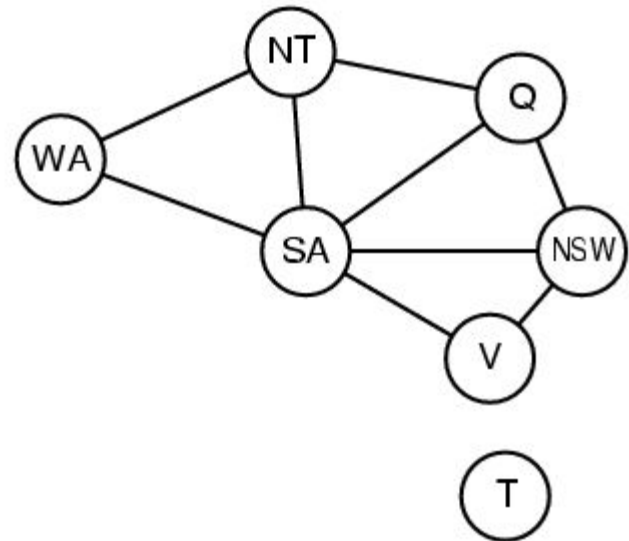
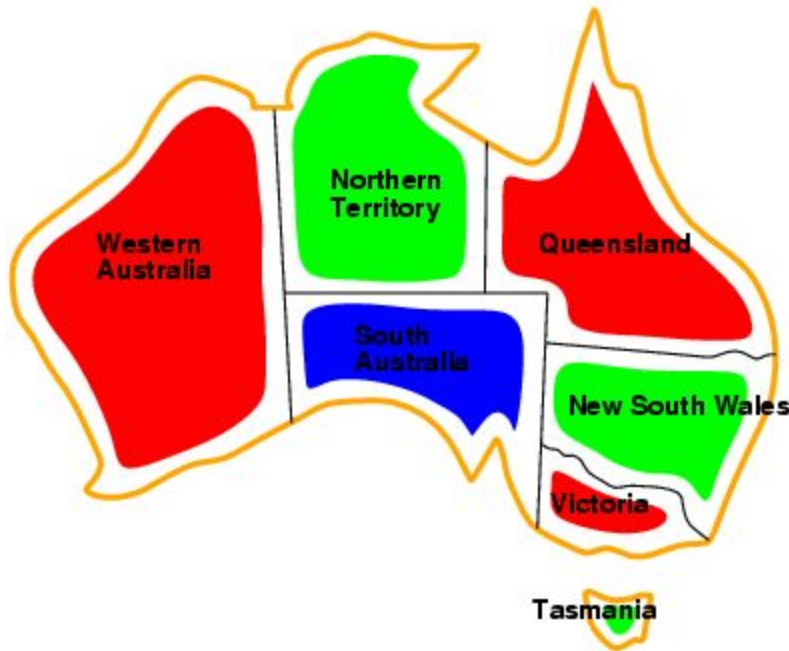
Example: Map-Coloring



- **Solutions** are **complete** and **consistent** assignments
- e.g., WA = red, NT = green, Q = red, NSW = green, V = red, SA = blue, T = green

Constraint graph

- **Binary CSP:** each constraint relates two variables
- **Constraint graph:** nodes are variables, arcs are constraints



Varieties of constraints

- **Unary** constraints involve a single variable,
 - e.g., $SA \neq \text{green}$
- **Binary** constraints involve pairs of variables,
 - e.g., $SA \neq WA$
- **Higher-order** constraints involve 3 or more variables,
 - e.g., cryptarithmic column constraints

Backtracking search

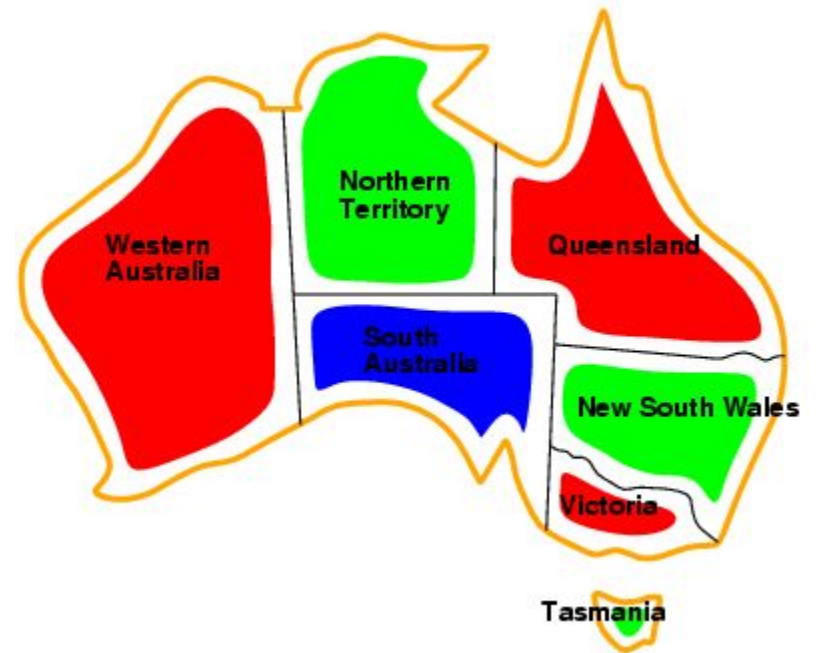
- Variable assignments are **commutative**, i.e.,
[WA = red then NT = green] same as [NT = green then WA = red]
- \Rightarrow Only need to consider assignments to a single variable at each node
- Depth-first search for CSPs with single-variable assignments is called backtracking search
- Can solve n -queens for $n \approx 25$

Backtracking search

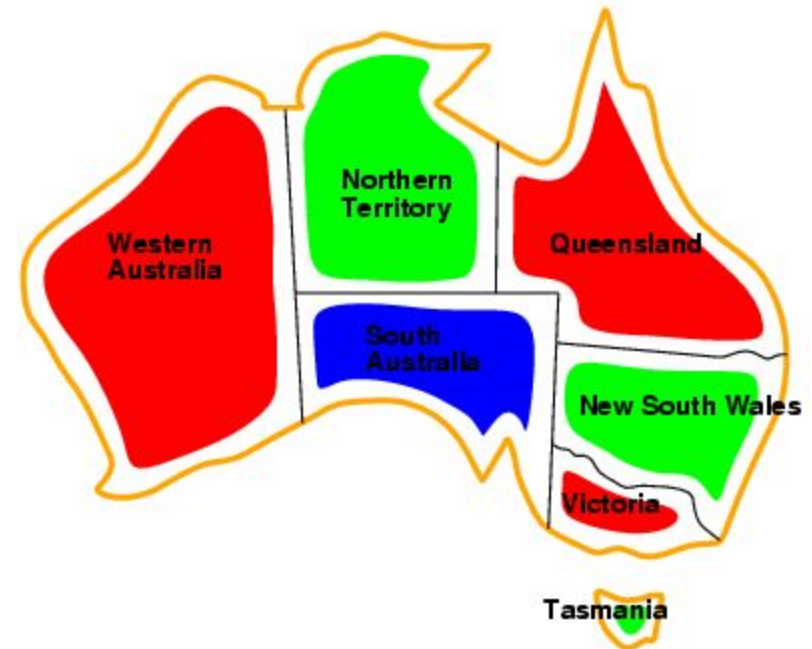
```
function BACKTRACKING-SEARCH(csp) returns a solution, or failure
  return RECURSIVE-BACKTRACKING({}, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns a solution, or failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(Variables[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment according to Constraints[csp] then
      add { var = value } to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove { var = value } from assignment
  return failure
```

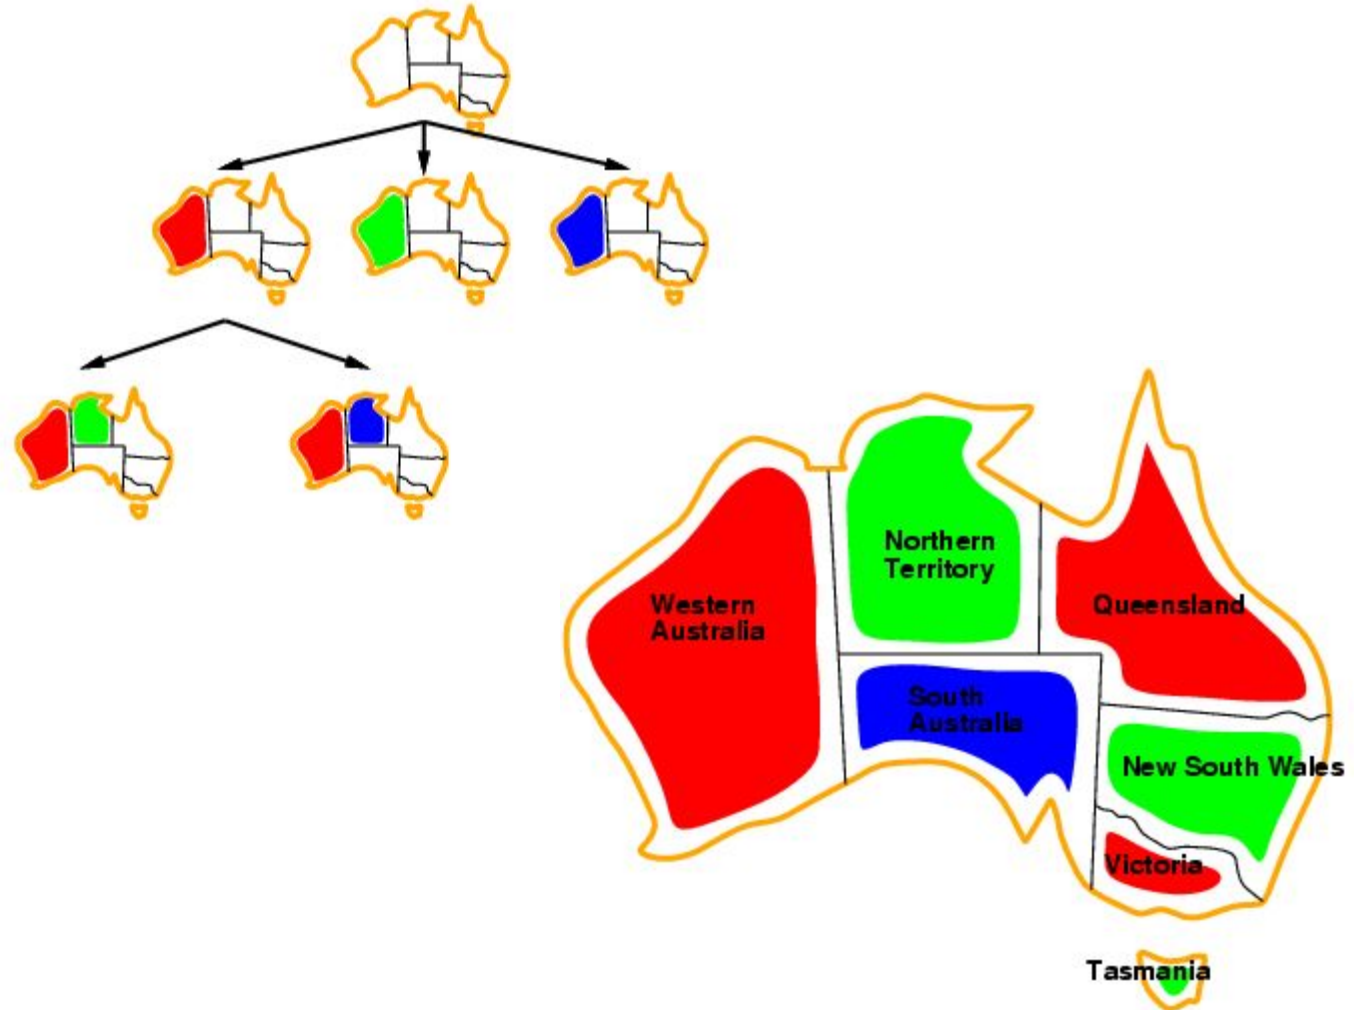
Backtracking example





Backtracking example



Backtracking example



Backtracking example



Improving backtracking efficiency

- **General-purpose** methods can give huge gains in speed:
 - Which variable should be assigned next?
 - In what order should its values be tried?
 - Can we detect inevitable failure early?

Variable and Value Selection


- Selecting variables and assigning values using a static list is not always the most efficient approach.
 - Difficult to make the "right" choice for picking and setting the next variable.

HEURISTICS can help here, e.g.,


- "Minimum remaining values" heuristic - choose the variable with the smallest number of remaining values in its domain.
 - Also called "most constrained variable" heuristic
- "Degree heuristic" - choose the variable that is part of the most remaining unsatisfied constraints.
 - Useful to select first variable to assign.
- "Least-constraining-value" heuristic - once a variable is chosen, choose its value as the one that rules out the fewest choices for neighboring variables.
 - Keeps maximum flexibility for future variable assignments.

Most constrained variable

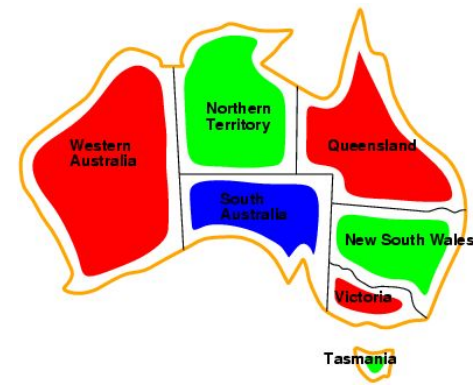
- Most constrained variable:
choose the variable with the fewest legal values




- a.k.a. **minimum remaining values (MRV)** heuristic



Degree Heuristic




- A good idea is to use it as a tie-breaker among most constrained variables
- Most constraining variable:
 - choose the variable with the most constraints on remaining variables



Least constraining value


- Given a variable to assign, choose the least constraining value:
 - the one that rules out the fewest values in the remaining variables



Forward checking

- Idea:

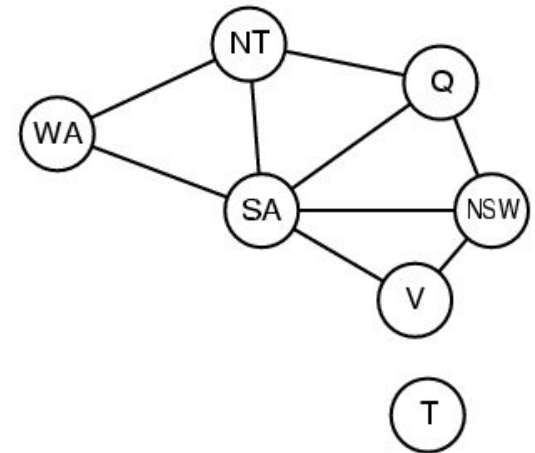
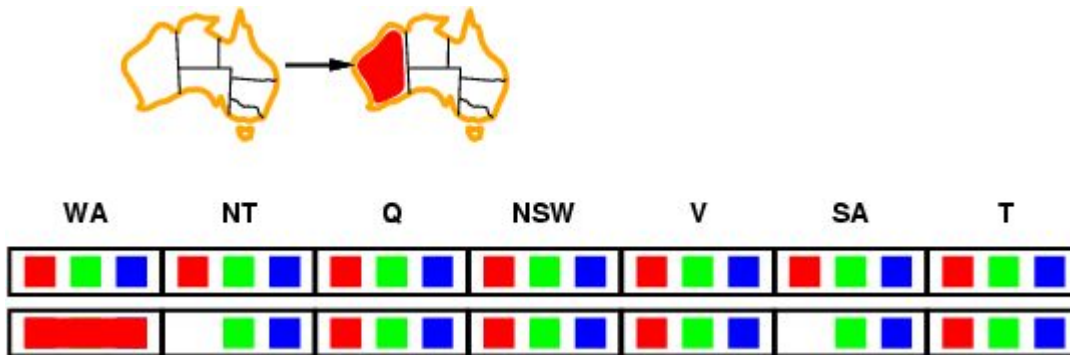
- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values



Forward checking

- Idea:



- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values



Forward checking

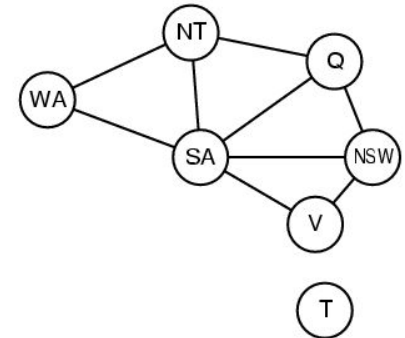
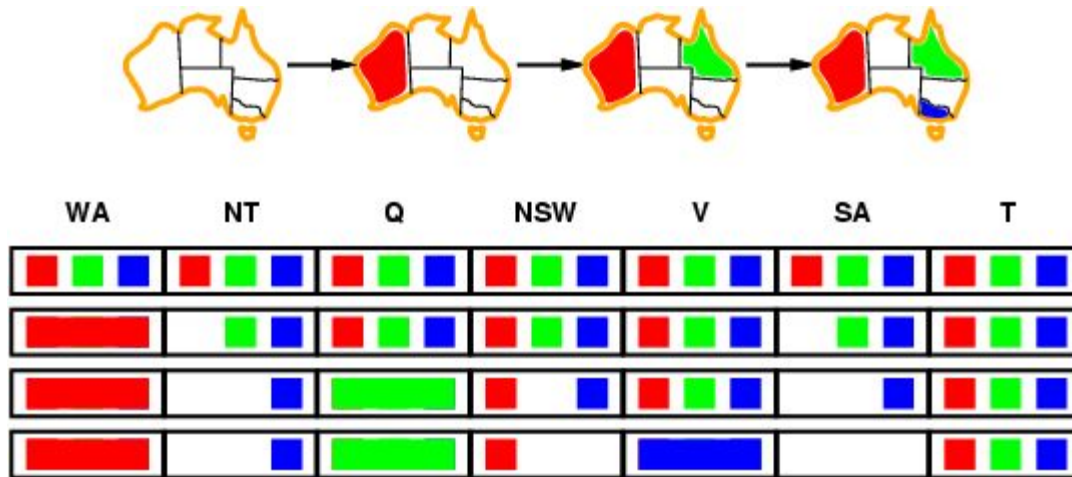
- Idea:

- Keep track of remaining legal values for unassigned variables
- Terminate search when any variable has no legal values



Forward checking


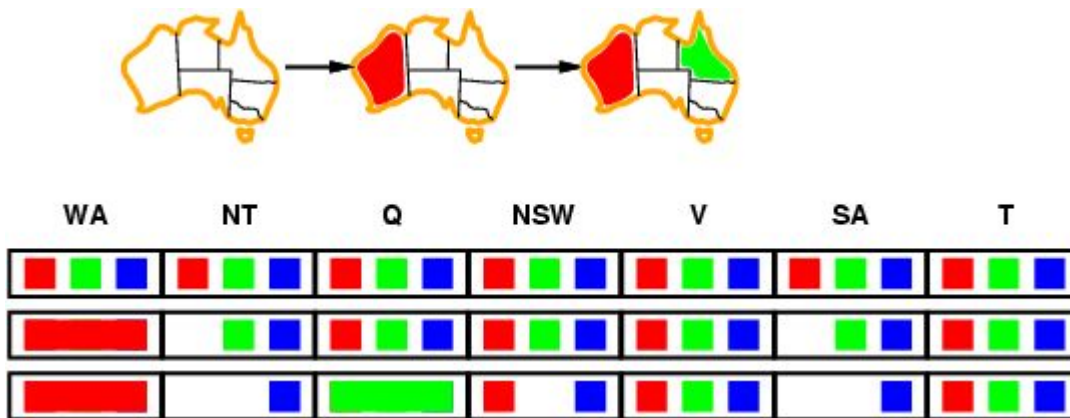
- Idea:
 - Keep track of remaining legal values for unassigned variables
 - Terminate search when any variable has no legal values



	WA	NT	Q	NSW	V	SA	T
Domains	RGB	RGB	RGB	RGB	RGB	RGB	RGB
After WA	(R)	GB	RGB	RGB	RGB	GB	RGB
After Q	(R)	B	(G)	RB	RGB	B	RGB
After V	(R)	B	(G)	R	(B)	 	RGB

Constraint propagation

- Forward checking propagates information from assigned to unassigned variables, but doesn't provide early detection for all failures:

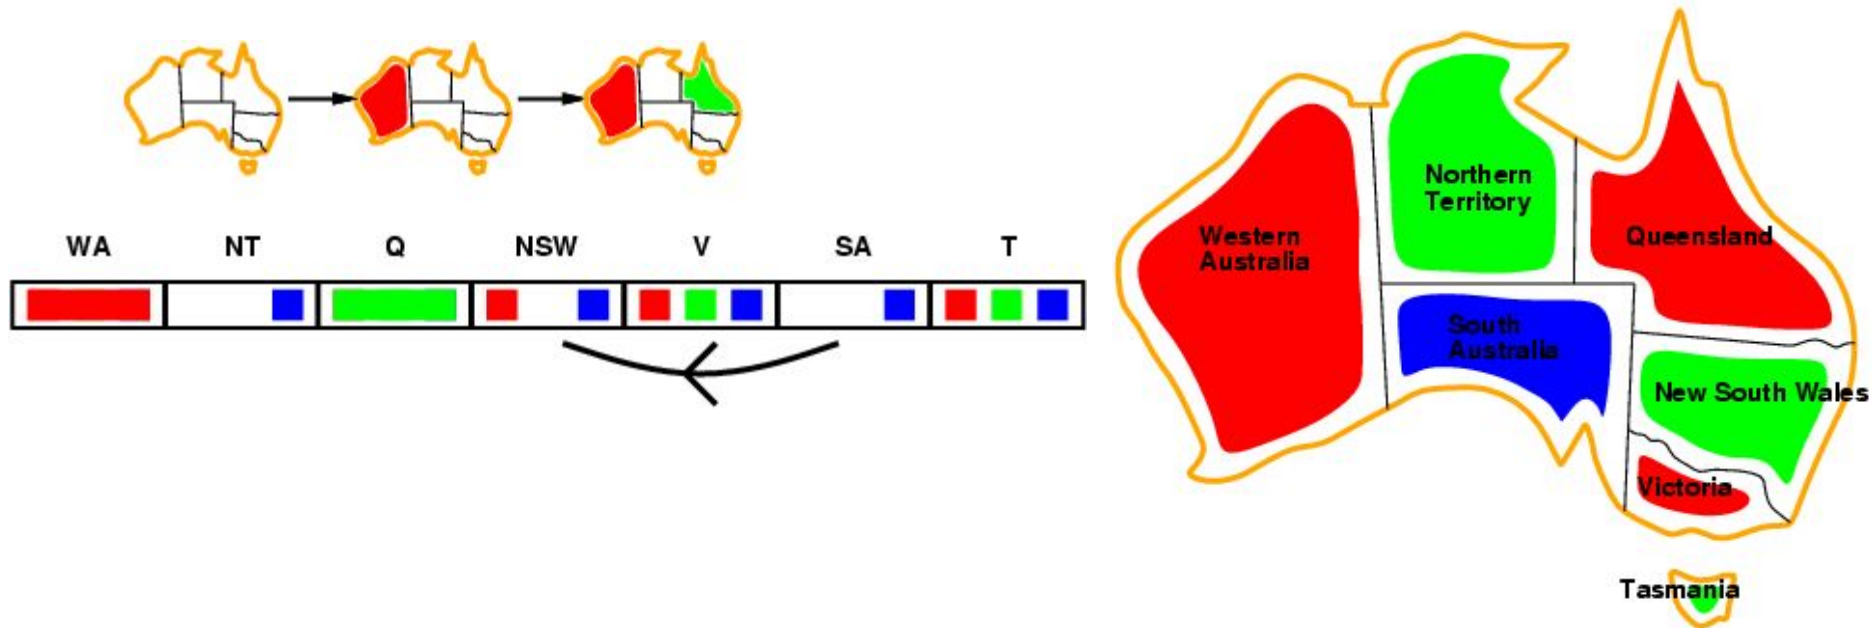


- NT and SA cannot both be blue!
- Constraint propagation algorithms repeatedly enforce constraints locally...

Arc consistency

- Simplest form of propagation makes each arc **consistent**
- $X \sqcap Y$ is consistent iff



for **every** value x of X there is **some** allowed y



Arc consistency

- Simplest form of propagation makes each arc **consistent**
- $X \square Y$ is consistent iff

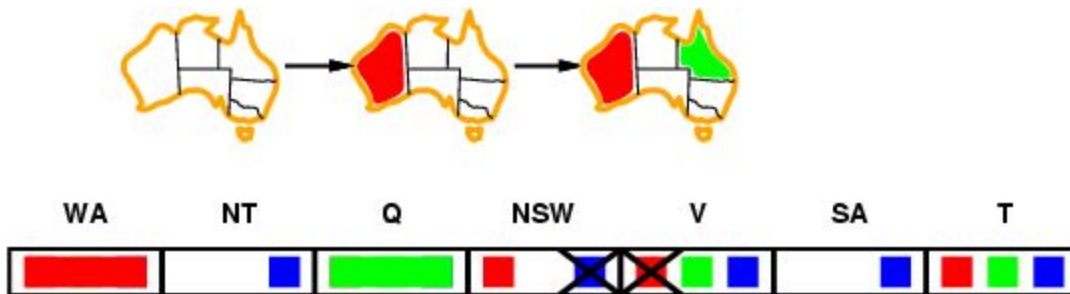
for **every** value x of X there is **some** allowed y




Arc consistency

- Simplest form of propagation makes each arc **consistent**
- $X \sqcap Y$ is consistent iff

for **every** value x of X there is **some** allowed y




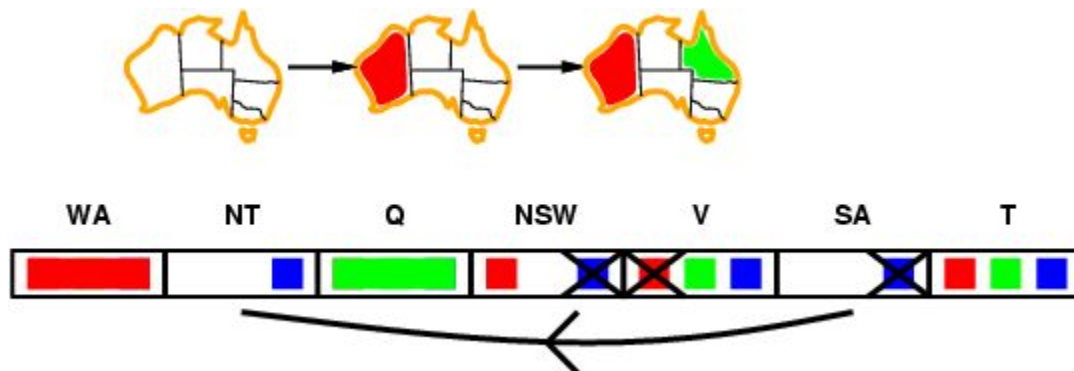
if X loses a value, neighbors of X need to be rechecked



Arc consistency

- Simplest form of propagation makes each arc **consistent**
- $X \sqcap Y$ is consistent iff

for **every** value x of X there is **some** allowed y



- If X loses a value, neighbors of X need to be rechecked
- Arc consistency detects failure earlier than forward checking
- Can be run as a preprocessor or after each assignment

Arc consistency algorithm AC-3

```
function AC-3(csp) returns the CSP, possibly with reduced domains
  inputs: csp, a binary CSP with variables  $\{X_1, X_2, \dots, X_n\}$ 
  local variables: queue, a queue of arcs, initially all the arcs in csp

  while queue is not empty do
     $(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\textit{queue})$ 
    if RM-INCONSISTENT-VALUES( $X_i, X_j$ ) then
      for each  $X_k$  in NEIGHBORS[ $X_i$ ] do
        add  $(X_k, X_i)$  to queue



---


function RM-INCONSISTENT-VALUES( $X_i, X_j$ ) returns true iff remove a value
  removed  $\leftarrow$  false
  for each  $x$  in DOMAIN[ $X_i$ ] do
    if no value  $y$  in DOMAIN[ $X_j$ ] allows  $(x, y)$  to satisfy constraint( $X_i, X_j$ )
      then delete  $x$  from DOMAIN[ $X_i$ ]; removed  $\leftarrow$  true
  return removed
```

- Time complexity: $O(\#\text{constraints} |\text{domain}|^3)$

Checking consistency of an arc is $O(|\text{domain}|^2)$